November 2026

Editor-in-Chief Journal of the Association for Information Science and Technology

Dear Editor,

I submit the manuscript “Which reverts count? Reverter populations, attribution rules and observation windows behind a Wikipedia indicator of information quality work” for consideration as a Research Article. It has not been published and is not under consideration elsewhere. I am the sole author.

Connection to information science and to this journal. The information quality framework of Stvilia et al. (2007), published in this journal, built two of its Wikipedia measures on reverts without saying how a revert was identified or whose reverts were counted. The manuscript starts from that observation and asks what the revert, now a standard trace of information quality work, is an indicator of. A review of 41 studies that use a revert measure, three of them from this journal and the Journal of Documentation, coded with language-model assistance and then checked by me against the source texts, indicates that it stands for four constructs (a quality judgment, conflict, a contributor’s experience of rejection, human effort) and that the choices each construct requires are often unreported. I report a conservative tally that also counts choices recoverable from a cited procedure, a library default or released code: 22 studies give no indication of how bots are handled, 28 no observation window, and none varies the reverter population. The review is a structured sample read by one person, and the coding table with the passage behind each code is released. The manuscript then measures those choices on 3.2 million newcomer first edits in seven language editions.

What is new. (1) MediaWiki’s own revert tag records a revert only when the reverting edit is trusted or patrolled. In the two editions that patrol recent changes it omits, for newcomer first edits, 71 and 79 percent of reverts made by unregistered editors and 60 to 70 percent of those by registered accounts without patrol-exempt rights, against almost none in editions without patrolling; the manuscript tests this directly with the user groups each reverting account held at the time. The consequence is concrete: the difference between English and Italian newcomer revert rates is 3.6 points under a hash-based rule and 7.4 points when measured with the tag, and the manuscript decomposes the change cell by cell. The tag’s conditions are documented in the

software manual; I found no published measurement of their extent, while tag-based indicators have begun to appear in the literature. (2) A revert rate is a mixture: bots, one-click tools used by a small number of accounts, other editors, unregistered editors and the newcomers themselves each contribute a measured share that differs across editions, and the newcomers reverted by each class differ by a factor of 1.4 to 4 in their odds of returning between tool and manual reverts. (3) Whether excluding a class of reverts matters less with a longer window depends on the attribution rule. (4) Proposed reporting items tie the choices to the construct an analyst intends. That reverters differ in kind was known for the English edition; the contribution here is measurement across editions and rules, and I have tried not to claim more.

Registration. Four editions were explored; hypotheses, thresholds and code were then registered on Zenodo (https://doi.org/10.5281/zenodo.23720117), and three held-out editions were opened afterwards. This is a hold-out replication, not a blind preregistration, and the manuscript says so. All four registered outcomes are reported, including a clause that was rejected and a hypothesis that held in two of three editions. The registered tests concern a filter that, as the review found, published studies do not use; I report them briefly and as registered. The findings listed above as (1) and (2), and the rule dependence in (3), are exploratory and are labeled as such in the abstract, tables and figures. A deviation log is public.

Generative AI. Generative AI (Claude, Anthropic) was used for three tasks: drafting analysis code, drafting and editing prose, and, in the review of practice, retrieving and screening texts, extracting passages and proposing codes. I checked the codes of all 41 studies against the sources; there was no second human coder, and the manuscript claims no reliability coefficient for the final codes. A second model (ChatGPT, OpenAI) assisted that check for four studies. Model agents were also used to critique drafts. No model generated, altered or imputed the audit's data; every number of the audit is produced by released deterministic code from unmodified public dumps. I specified the design, made the analytic decisions, checked the code and the review coding against the sources, and take responsibility for every statement. The same declaration appears in the manuscript, and a dated log of AI assistance is in the repository.

Data, code and related work. The data are public Wikimedia dumps. Code, aggregates, the review table and logs are on GitHub (https://github.com/beatsbywoni/wiki-revert-audit) and archived on Zenodo (https://doi.org/10.5281/zenodo.24054641). A separate manuscript of mine

applies a related audit logic to software repository mining; it uses different data and shares no results with this one.

The body is about 7,300 words including five tables, figure captions and endnotes (about 6,800 without the tables), with a 200-word abstract; the Supplementary Material holds the registration account, eleven tables and two figures. I have no funding and no conflicts of interest to declare. Suggested reviewers are entered in the submission system.

Thank you for considering the manuscript.

Sincerely,

Hongwon Jeong Center for Creative Convergence Education, Hanyang University ERICA, Ansan, Republic of Korea ORCID 0009-0003-7908-7478 beatsbywoni@hanyang.ac.kr