November 2026

Editor-in-Chief Journal of the Association for Information Science and Technology

Dear Editor,

I submit the manuscript “Which reverts count? Reverter populations, attribution rules and observation windows behind a Wikipedia indicator of information quality work” for consideration as a Research Article. It has not been published and is not under consideration elsewhere. I am the sole author.

Connection to information science and to this journal. The information quality framework of Stvilia et al. (2007), published in this journal, built two of its Wikipedia measures on reverts without saying how a revert was identified or whose reverts were counted. The manuscript starts from that observation and asks what the revert, now a standard trace of information quality work (Stvilia et al., 2008), is an indicator of. It reads the trace as a record of one act of quality work that is used for four constructs (a quality judgment, conflict, a contributor’s experience of rejection, human effort), and it measures what the record contains on 3.2 million newcomer first edits in seven language editions. The manuscript is an empirical audit, not a literature review: a survey of reporting practice in 41 studies motivates it and is presented as such, with its limits (one reader, language-model assistance, no second human coder).

What is new. (1) The platform’s own revert tag is approval-conditioned documentation. MediaWiki applies the tag only when the reverting edit meets the software’s approval condition, at once if the reverter is exempt from patrol and otherwise only if someone patrols the edit. In the Italian edition, which patrols recent changes, 79 percent of reverts of newcomer first edits made by unregistered editors and 77 percent of manual reverts by registered editors without a patrol-exempt group carry no tag, against about 1 percent for editors with a recorded patrol-exempt group; in French the omission follows the edition’s automatic promotion at 500 edits. The condition is documented in the software manual; I found no published measurement of its extent. Separately, replacing a hash-based rule by the tag changes the difference between English and Italian newcomer revert rates from 3.6 to 7.4 points, and the manuscript decomposes the change. For information science this is a finding about documentation and invisible work (Star & Strauss, 1999): for the reverts of newcomer edits in the editions studied, whether an act of quality work is

captured by the tag varies with the standing of the worker, in the pattern the documented approval rule predicts. The manuscript reports this as an association with recorded user groups and edit counts; patrol logs were not joined, and the revision history itself records every revert. (2) A revert rate is a mixture: bots, one-click tools used by a small number of accounts, other editors, unregistered editors and the newcomers themselves each contribute a measured share that differs across editions, and the newcomers reverted with tools and manually differ by a factor of 1.4 to 4 in their odds of returning. (3) Proposed reporting items tie the choices to the construct an analyst intends. That reverters differ in kind was known for the English edition; the contribution here is measurement across editions and rules, and I have tried not to claim more.

Registration. Four editions were explored; hypotheses, thresholds and code were then registered on Zenodo (https://doi.org/10.5281/zenodo.23720117), and three held-out editions were opened afterwards. This is a hold-out replication, not a blind preregistration, and the manuscript says so. All four registered outcomes are reported, including a clause that was rejected and a hypothesis that held in two of three editions. The registered tests concern a filter that the surveyed studies do not use; I report them briefly and as registered. The findings listed above as (1) and (2) are exploratory and are labeled as such in the abstract, tables and figures. A deviation log is public.

Generative AI. Generative AI (Claude, Anthropic) was used for three tasks: drafting analysis code, drafting and editing prose, and, in the survey of reporting practice, retrieving and screening texts, extracting passages and proposing codes. I checked the codes of all 41 studies against the sources; there was no second human coder, and the manuscript claims no reliability coefficient. A second model (ChatGPT, OpenAI) assisted that check for four studies, two of them among the 41. Model agents were also used to critique drafts, and several analyses were added in response. No model generated or altered the audit's data, and none analyzed them on its own; every number of the audit is produced by released deterministic code from unmodified public dumps. I specified the design, made the analytic decisions, checked the code and the review coding against the sources, and take responsibility for every statement. The same declaration appears in the manuscript, and a dated log of AI assistance is in the repository.

Data, code and related work. The data are public Wikimedia dumps. Code, aggregates, the coding table of the survey and logs are on GitHub (https://github.com/beatsbywoni/wiki-revert-audit) and archived on Zenodo (https://doi.org/10.5281/zenodo.24054641). A separate

manuscript of mine applies a related audit logic to software repository mining; it uses different data and shares no results with this one.

The manuscript has about 6,900 words including the body, four tables, captions and endnotes, and excluding the abstract, keywords and references, with a 200-word abstract; the Supplementary Material holds the registration account, eleven tables and three figures. I have no funding and no conflicts of interest to declare. Suggested reviewers are entered in the submission system.

Thank you for considering the manuscript.

Sincerely,

Hongwon Jeong Center for Creative Convergence Education, Hanyang University ERICA, Ansan, Republic of Korea ORCID 0009-0003-7908-7478 beatsbywoni@hanyang.ac.kr